
How to Pick the Model: A Framework for Budget-Aware Orchestration

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Agent systems on verified tasks invite a deceptively simple question: which back-
2 bone model should be run? The strongest model is not always the best deployment
3 choice, because verified progress must be bought under finite wall-clock, token,
4 API, and evaluation budgets. We formulate this problem as deployment-aware
5 model selection. A fixed router observes a pre-deployment information state and
6 chooses a worker or harness action; the decision is scored by a checker-indexed
7 deployment loss rather than by routing accuracy or final checker score alone. The
8 key estimand is paired: on the same task instances, with the same router and action
9 menu, does richer information change the selected action and reduce net deploy-
10 ment loss after its own measurement and context costs are charged? We instantiate
11 the framework on a Karpathy AutoResearch-inspired CIFAR-10 edit-verify bench-
12 mark, where a prompt-only router chooses among heterogeneous workers before
13 costly deployment runs. The protocol reports when checker-score rankings disagree
14 with deployment-loss rankings, when pre-deployment information changes routing,
15 and whether those changes pay for themselves. The resulting protocol evaluates
16 not which agent is best in isolation, but when information is worth buying before
17 deploying one.

18 1 Introduction

19 Practitioners building agent systems face multiple practical challenges: ~~what model~~which model
20 (backbone) to use, how to route between models, when to spawn sub-agents, how to engineer context
21 between them, and so on. All these possible choices often affect performance (quality of the solution),
22 speed (time to solution), and ~~the mere token cost~~token/API cost. The question is how to ~~pick between~~
23 ~~the different possible designs~~choose among the possible designs. The common way to answer this
24 question is benchmark racing: take a benchmark ~~as~~such as SWE-bench [1] or MLE-bench [2]
25 and pick the agent system with the best average end-to-end performance. An increasingly long line
26 of work ~~, based on this “recipe”,~~based on this recipe is proposing better and better ~~agentie-system~~
27 ~~arehitecture~~agentie system architectures to achieve ~~SOTA at~~state-of-the-art performance on these
28 tasks. ~~This article investigates:~~This article investigates

29 ~~whether there exist other ways to answer the question of what the most suited agent system is:~~
30 ~~whether there are other ways to answer the question of which agent system is most suitable.~~

31 Indeed, (i) ~~a large number of tasks for which such benchmark does not exist (as developing academic~~
32 ~~research [?]).~~ And (ii) ~~very often, the task presents instances of varied difficulty~~there are many tasks
33 ~~for which such benchmarks do not exist, such as developing academic research [3].~~ Moreover, tasks
34 ~~often contain instances of varied difficulty~~and the practitioner may want to understand whether, ~~on~~
35 ~~the single instance~~on a single instance, a lower-cost, faster model is sufficient; ~~more tools should~~

36 ~~be given to the system~~the system should be given more tools; or a more involved orchestration is
 37 required. As a running example for the rest of the paper, ~~the question whether the user should use~~
 38 ~~on their own task Haiku vs Sonnet vs Opus in Claude Code, or Claude Code vs Codex,~~the question
 39 of whether a user should use Haiku, Sonnet, or Opus in Claude Code, or Claude Code rather than
 40 Codex, for their own task is not fully addressed by benchmark racing.

41 The central issue is that one benchmark score may ~~conflates~~conflate mechanisms that imply dif-
 42 ferent engineering decisions. A system can win because its worker ~~takes~~has lower per-step time
 43 or ~~token count~~. ~~Because the selected worker is~~token cost, or because the selected worker is more
 44 competent on a solver-relevant task mode. Analogously, imagine ~~some sort of pre-deployment~~
 45 ~~signal~~a pre-deployment signal (~~problem definition, the context, or an initial test~~the problem defi-
 46 nition, context, or an initial test) reveals information about the instance ~~which~~that can be used ~~to~~
 47 ~~better orchestrate or to choose what backbone to pick~~to choose a better backbone or orchestration.
 48 A system can win ~~if it can orchestrate or choose the model optimally exploiting that information or~~
 49 ~~not~~by exploiting that information when choosing a model or orchestration, or it can lose by failing
 50 to exploit it—as humans implicitly do when they decide to use lower or higher reasoning effort
 51 ~~reasoning~~ or different models on their tasks.

52 More practically, in this article *we study the deployment problem of choosing which model or harness*
 53 *to run on a verifiable task*. The objective is ~~time to a required verified result~~time to reach a required
 54 ~~verified result~~: among actions that may eventually succeed, the useful one is the one expected to
 55 reach the verifier threshold fastest, after charging any pre-deployment measurement used to make
 56 the choice. ~~Achille and Soatto, indeed, recently showed that for verifiable tasks the time to solution,~~
 57 ~~as in their Proper Time, is the performance measure to compare them.~~ Indeed, Achille and Soatto
 58 [4] recently argued that time to solution, as formalized by their Proper Time, is a performance
 59 measure for comparing agents on verifiable tasks. ~~Building on this intuition we operationalize the~~
 60 ~~difference in time to solution of two possible designs. We show it mathematically decomposes in 4~~
 61 ~~components which isolates the mechanisms of choice mentioned above.~~ Building on this intuition,
 62 we operationalize a certified log-effort surrogate gap between two possible designs. We show that it
 63 mathematically decomposes into four components that isolate the mechanisms of choice mentioned
 64 above. ~~The experimental protocol then shows these can be estimated and lead to improvements~~
 65 ~~in speed on Karpathy’s AutoResearch tasks.~~ The experimental protocol estimates the operational
 66 quantities that are directly measurable and tests, by paired selected-worker deployment, whether richer
 67 pre-deployment information improves net deployment loss after measurement costs are charged.

68 **Research questions and contributions.** We organize the paper around three questions, each paired
 69 with the contribution that answers it.

- 70 1. Can the ~~time-to-solution~~certified log-effort surrogate gap between agent designs be decom-
 71 posed into meaningful terms? Theorem 1 gives an exact four-term identity under a packed
 72 latent-mode assumption: cost, within-mode competence, routing information, and a predeclared
 73 mode-summary divergence additively explain the improvement in expected certified log-effort.
- 74 2. Can the decomposition ~~replace~~structure or ~~reduce~~ end-to-end racing for model selection?
 75 Algorithm 1 and the pilot concentration bound in Appendix A turn the identity into a structured
 76 pilot protocol whose evaluation cost depends on the number of models and task modes.
- 77 3. When does a lower-cost model with retries beat a stronger model? Theorem 2 gives the
 78 single-mode crossover rule, and the multi-mode accounting shows why the answer changes across
 79 task modes.

80 The experiments below instantiate these claims with checked trajectories and deployment accounting.

81 2 Problem setup and deployment estimands

82 We study the problem faced by a deployment system before launching an agentic run: given a task
 83 instance and a finite menu of possible workers or harnesses, which one should be run? The answer
 84 is not necessarily the model with the best final score. A deployment action must produce checked
 85 progress under finite time, token, API, and evaluation budgets.

86 **Definition 1** (Budgeted externally checked task). A ~~checked-deployment task~~ *budgeted externally*
 87 *checked task* is a tuple

$$\mathcal{T} = (\mathcal{X}, \mathcal{Y}, V, B_{\text{dep}}, \mu).$$

88 Here $\mathcal{X}$ is an instance space, $\mathcal{Y}$ is a ~~artifact-or-solution space~~ *solution space*, $V : \mathcal{X} \times \mathcal{Y} \rightarrow \mathcal{R}$ is an
 89 external checker or scoring procedure on instance–artifact pairs, B_{dep} is the full deployment budget
 90 and acceptance object, and μ is the deployment distribution over instances. The object B_{dep} may
 91 include a horizon, verifier budget, token/time/cost caps, and a thresholding rule that induces a binary
 92 success event from the checker score. For a realized instance $x \in \mathcal{X}$, an agent produces artifacts
 93 $y \in \mathcal{Y}$ that are scored by $V(x, y)$. An agent succeeds on $x \sim \mu$ if it produces an accepted y before
 94 exhausting B_{dep} .

95 The checker may be binary, as in unit tests or proof checking, or scalar, as in validation loss or reward.
 96 The key point is externality: progress is determined by the checker, not by the model’s self-report.

97 **Definition 2** (Workers, actions, and routers). Let $\mathcal{M} = \{M_1, \dots, M_K\}$ be a finite menu of workers,
 98 where a worker is any deployable solver or fixed agent policy that can produce a candidate solution
 99 for verification. Let $\mathcal{A}$ be a finite menu of deployable actions. An action may be a single model, a
 100 fixed model harness, a retry policy, or a tool-using agent protocol. In the simplest setting, $\mathcal{A} = \mathcal{M}$;
 101 more generally, an action bundles a worker with fixed stopping, retry, tool-use, and carryover rules.
 102 ~~The deployment loss is defined locally in this action definition.~~ The deployment loss is defined in
 103 Section 4, where it is tied to first-passagage cost, audit loss, and measurement loss. A routed design is a
 104 tuple

$$d = (\mathcal{A}, R, \kappa),$$

105 where R is a router that selects an action after observing permitted information and κ is the declared
 106 scalar action-cost convention used in the log-effort diagnostic. When costs are logged as resource
 107 vectors, the protocol must fix the scalar reduction before evaluation. A router R observes a permitted
 108 pre-deployment information state ~~$W(x)$~~ $Z_j(x)$ and selects an action

$$a_j(x) = R(Z_j(x)) \in \mathcal{A}.$$

109 *Equivalently, a signal condition induces the routing policy $\rho_j : x \mapsto a_j(x)$. The selected action then*
 110 *performs the deployment run from the original task state.*

111 This separates routing from deployment. The router chooses what to run; the selected action performs
 112 the actual checked run. Unless explicitly declared otherwise, pre-deployment records are shown only
 113 to the router and are not passed to the selected worker.

114 **Definition 3** (Task instances and task modes). Given $\mathcal{T} = (\mathcal{X}, \mathcal{Y}, V, B_{\text{dep}}, \mu)$, a task instance is
 115 a concrete realization $X = x \in \mathcal{X}$ of that same task. It is the measurable object supplied to the
 116 deployed system; the router and worker may access only the fields made available by the protocol. A
 117 task mode is a latent variable

$$S = \sigma(X) \in \mathcal{S} = \{1, \dots, m_S\}$$

118 induced by a measurable map $\sigma : \mathcal{X} \rightarrow \mathcal{S}$ that assigns each instance to one of finitely many solver-
 119 relevant variants of the task. It defines a prior $\pi_s = \mathbb{P}_\mu[S = s]$ and mode-conditional instance
 120 distributions $\mu_s = \mu(\cdot \mid S = s)$. Modes are therefore not different tasks and not worker labels. They
 121 are variants of the same task: finite subpopulations of instances that share solver-relevant structure
 122 and may favor different workers or interventions.

123 The mode variable is retained by the evaluator for accounting and diagnostic comparisons. It need not
 124 be observed by either the router or the worker. This lets the same theory cover expert-defined task-
 125 mode schemas and finite evaluator-defined modes. When an experiment uses one fixed executable
 126 prototype per mode, it estimates conditional panel quantities rather than the population quantities
 127 induced by μ .

128 **Assumption 1** (Packed latent-mode family). Each instance $X \sim \mu$ is associated with a mode
 129 $S = \sigma(X)$ in a finite mode set $\mathcal{S}$, with prior π . Conditional on mode s , action a has success
 130 probability $p(a, s) > 0$ and scalar cost

$$\kappa(a, s) = \mathbb{E}[C_\delta(a, X; \xi) \mid S = s] > 0$$

131 under the declared deployment protocol, where C_δ is the normalized cost accumulated to the first
 132 certified success or the horizon. At deployment time, the router observes a signal $Z = z$ and selects
 133 one action. The evaluator defines the mode distribution induced by that signal,

$$\pi_z(s) = \mathbb{P}[S = s \mid Z = z].$$

134 The router is not required to output π_z or any mode distribution. Instead, the theory uses a fixed
 135 evaluator-side mode-summary distribution $q_z \in \Delta(\mathcal{S})$, with $q_z(s) > 0$ whenever $\pi_z(s) > 0$. This
 136 summary is not a router output; it is a theory-facing diagnostic distribution specified before evaluation.
 137 Write $a_z = \rho(z)$ for the action selected after observing z . For an instance whose true mode is s , the
 138 certified effort surrogate is proportional to

$$T_{\rho,q}(s, z) \propto \frac{\kappa(a_z, s)}{p(a_z, s) q_z(s)}. \quad (1)$$

139 Assumption 1 is deliberately restrictive. Real task instances can vary in infinitely many ways, while
 140 the theory uses a finite state abstraction. The assumption says that, for model selection, instances
 141 can be packed into modes with stable action success probabilities. The purpose is to isolate the
 142 regime in which cost, competence, information, and the predeclared mode-summary divergence are
 143 algebraically separable. Residual and calibration diagnostics measure when this finite packing is too
 144 coarse.

145 **Definition 4** (Cost-adjusted certified log-effort diagnostic objective). *For a fixed routed policy*
 146 $\rho : Z \mapsto a_Z$ *and a fixed evaluator summary* q *satisfying Assumption 1, define*

$$\mathcal{J}(\rho, q) = \mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} [\log \kappa(a_Z, S) - \log p(a_Z, S)] + \mathbb{E}_Z [\mathcal{H}(\pi_Z) + \text{KL}(\pi_Z \| q_Z)]. \quad (2)$$

147 Lower $\mathcal{J}$ means lower expected certified log-effort under the fixed diagnostic convention. The quantity
 148 q is not optimized by the router; it is fixed by the evaluator-side diagnostic rule. When several signal
 149 conditions are compared, $\mathcal{J}$ is applied to the corresponding policy summaries; Appendix C writes
 150 this comparison explicitly for Z_j versus Z_0 .

151 The action quantities in (2) are $\kappa(a_Z, s)$ and $p(a_Z, s)$: how costly the selected action is on mode s
 152 and how often it succeeds there. The router and information quantities are Z , π_z , and q_z : what the
 153 router is allowed to observe, what that signal reveals about the task mode, and the evaluator-side
 154 mode-summary distribution used in the KL term.

155 **Remark 1** (Meaning of the effort surrogate). *Equation (1) says that, on a true mode s after signal z ,*
 156 *certified effort increases with action cost, decreases with action success probability on that mode,*
 157 *and increases when the evaluator-side mode summary gives too little diagnostic mass to that mode.*
 158 *The router is not literally routing jobs to modes; modes are evaluator-defined task variants. The*
 159 *quantity $q_z(s)$ is a diagnostic summary, not a deployment probability. The KL term below is therefore*
 160 *a mode-summary divergence under the chosen diagnostic convention. It should be interpreted as a*
 161 *router failure only if the experiment predefines a reproducible construction of q_z from router behavior.*

162 3 Performance accounting

163 For a fixed signal value z , the routing part of the objective is the cross-entropy between the signal-
 164 induced mode distribution and the fixed evaluator-side mode-summary distribution:

$$\text{CE}(\pi_z, q_z) = - \sum_{s \in \mathcal{S}} \pi_z(s) \log q_z(s) = \mathcal{H}(\pi_z) + \text{KL}(\pi_z \| q_z). \quad (3)$$

165 The KL term is a divergence over finite task-mode distributions. It is not a token-level language-model
 166 divergence.

167 **Theorem 1** (Four-term routed-policy decomposition). *Under Assumption 1, compare a baseline*
 168 *policy ρ_0 that deploys fixed action a_0 without using Z and uses $q_0(\cdot | z) \equiv \pi$, with a deployed routed*
 169 *policy summarized by (ρ, q) . The improvement in expected certified log-effort, $\Delta = \mathcal{J}(\rho_0, q_0) -$
 170 $\mathcal{J}(\rho, q)$, *decomposes as**

$$\Delta = \underbrace{\mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} \left[\log \frac{\kappa(a_0, S)}{\kappa(a_Z, S)} \right]}_{\text{cost}} + \underbrace{\mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} \left[\log \frac{p(a_Z, S)}{p(a_0, S)} \right]}_{\text{competence}} + \underbrace{\mathbb{I}_\pi(S; Z)}_{\text{routing information}} - \underbrace{\mathbb{E}_Z [\text{KL}(\pi_Z \| q_Z)]}_{\text{mode-summary divergence}}. \quad (4)$$

171 *Proof sketch.* Taking logs in (1) gives $\log T_{\rho,q}(s, z) = C + \log \kappa(a_z, s) - \log p(a_z, s) - \log q_z(s)$.
 172 Conditioning on $Z = z$ and averaging over $S \sim \pi_z$ converts $-\mathbb{E} \log q_z(S)$ into the cross-entropy
 173 in (3). Subtracting the prior-matched baseline cancels constants and uses $\mathcal{H}(\pi) - \mathbb{E}_Z \mathcal{H}(\pi_Z) =$
 174 $\mathbb{I}_\pi(S; Z)$. The full algebra is given in Appendix A. $\square$

The four terms have different operational meanings. If the cost term is large and positive, a lower-cost action may be preferable even when it reaches success later in step count. If the competence term varies by mode, a global worker ranking is likely to hide comparative advantage. If information gain is high, the pre-deployment signal exposes solver-relevant structure. If the mode-summary divergence is high, the signal-induced mode distribution is poorly matched to the evaluator’s predeclared mode summary. For two signal conditions, the identity compares the routed policies induced by those signals. A richer signal is useful only if its induced action change improves the cost–success tradeoff after the cost of obtaining the signal is charged.

Operational reading. The task mode S is not an oracle label for the best worker. The evaluator fixes a mode schema before evaluation, then estimates the mode-conditional quantities $p(a, s)$ and $\kappa(a, s)$. Routing information asks whether a signal makes those modes more identifiable. Routing shift asks whether the selected action changes when the signal is enriched. Mode-summary divergence is the KL term defined by the fixed evaluator-side mode summary q . Allocation regret, used later in the empirical protocol, is a separate action-frontier diagnostic rather than an estimator of this KL term. The final decision is not made by the information term alone; it is made by paired deployment gain on held-out instances after measurement and context costs are charged. Appendix B gives the full operationalization of these diagnostics.

3.1 When lower-cost retries win

Theorem 2 (Single-mode retry crossover). *Fix a single homogeneous mode, success threshold δ , and horizon H . Let actions h and ℓ have one-attempt first-passage success probabilities $p_h = \mathbb{P}[\tau_\delta(h, X) \leq H]$ and $p_\ell = \mathbb{P}[\tau_\delta(\ell, X) \leq H]$, with expected first-passage costs κ_h and κ_ℓ . Assume $0 < p_\ell < p_h < 1$ and $\kappa_\ell < \kappa_h$. Assume retries of ℓ are conditionally i.i.d. fresh attempts, each with the same one-attempt horizon H and common Bernoulli success probability p_ℓ . The lower-cost retry depth that matches the strong action’s one-attempt success probability is*

$$D_{\text{cross}} = \frac{\log(1 - p_h)}{\log(1 - p_\ell)}.$$

The corresponding integer fixed-depth retry block is cost-favorable whenever

$$\lceil D_{\text{cross}} \rceil \kappa_\ell < \kappa_h.$$

Remark 2 (Stopping after first retry success). *If retries of the lower-cost action stop after the first success and each attempted retry has unconditional expected cost κ_ℓ , the expected cost at depth D is*

$$\kappa_\ell \sum_{d=1}^D (1 - p_\ell)^{d-1} = \kappa_\ell \frac{1 - (1 - p_\ell)^D}{p_\ell},$$

and the decision should also account for the residual failure probability under the deployment loss. Any difference in within-attempt time-to-success is already absorbed into κ_h and κ_ℓ .

The multi-mode case is where the four-term identity matters. A lower-cost worker can be rational on one task mode and irrational on another. A router is useful only when different actions win the cost–success tradeoff on different modes and the pre-deployment information helps identify which mode applies.

4 Decomposition-guided selection

The decomposition suggests a practical protocol. Rather than evaluate every complete design as an independent arm, the builder first estimates shared quantities: mode-conditional success, cost, predeclared signal diagnostics, and action-allocation diagnostics. Then a larger finite design menu can be scored analytically, with expensive deployment evaluations reserved for finalists and for paired validation.

The log-effort objective $\mathcal{J}$ is a diagnostic surrogate, not the deployment loss. For deployment validation we use a predeclared first-passage loss. Given a checked progress trace $G_1, \dots, G_H$, define

$$\tau_\delta(a, x; \xi) = \inf\{t \leq H : G_t(a, x; \xi) \geq \delta\}, \quad F_\delta(a, x; \xi) = \mathbf{1}\{\tau_\delta(a, x; \xi) = \infty\},$$

217 and

$$\text{Occ}_{\delta,H}(a, x; \xi) = H^{-1} \sum_{t=1}^H \mathbf{1}\{G_t(a, x; \xi) \geq \delta\}.$$

218 Let $c_\delta(a, x; \xi)$ be the normalized worker-side resource vector accumulated to $\tau_\delta \wedge H$, and let

$$C_\delta(a, x; \xi) = \lambda^\top c_\delta(a, x; \xi)$$

219 be the scalar cost used in $\kappa(a, s)$. The primary loss is

$$\ell_{\text{dep}}(a, x; \xi) = \alpha_{\text{fail}} F_\delta(a, x; \xi) + C_\delta(a, x; \xi), \quad (5)$$

220 with $\alpha_{\text{fail}} = 1$ and normalized resource weights λ fixed before evaluation. This is the loss used
 221 in $\mathcal{J}_R(j)$ and $\Delta_R(j)$ unless a different primary loss is explicitly declared. Because the benchmark
 222 records full trajectories to horizon H , we also report the full-horizon audit loss

$$\ell_{\text{audit}}(a, x; \xi) = \ell_{\text{dep}}(a, x; \xi) + \alpha_{\text{occ}}(1 - \text{Occ}_{\delta,H}(a, x; \xi)) + \alpha_{\text{qual}}\psi(G_H(a, x; \xi)), \quad (6)$$

223 where $\alpha_{\text{occ}} = \alpha_{\text{qual}} = 1$ and $\psi(u) = 1 - u$ after clipping $G_H \in [0, 1]$. The audit loss checks
 224 persistence and final quality; it does not change the first-passage cost convention in Theorem 1.

225 **Algorithm 1** (Theory-to-deployment selection protocol).

Input: Task family, checker, finite task-mode schema $\mathcal{S}$, action menu $\mathcal{A}$, admissible progressive signals $\{Z_j\}$, pilot cell count n_{cell} , deployment repeats Q_{dep} , confirmation budget.

Output: Recommended signal-conditioned deployment policy, reported accounting terms, paired deployment gain, and calibration diagnostics.

- 1 *Freeze the modes and loss.* Declare $\mathcal{S}$, success threshold δ , horizon H , cost normalizers, deployment loss (5), measurement costs, and router output contract.
- 2 *Structured pilot.* For each action $a \in \mathcal{A}$ and mode $s \in \mathcal{S}$, estimate $\hat{p}(a, s)$, $\hat{\kappa}(a, s)$, first-hit curves, occupancy, and uncertainty intervals.
- 226 3 *Signal and router diagnostics.* For each progressive signal Z_j , report any predeclared signal-information diagnostic and log the action allocation induced by the router.
- 4 *Accounting score.* Compute the four log-effort terms from Theorem 1 only for quantities whose diagnostic estimators were predeclared; otherwise report allocation regret separately.
- 5 *Paired deployment validation.* On the same holdout instances, run the selected action under Z_0 and richer Z_j conditions, score repeated selected-worker trajectories by (5), and subtract measurement cost.
- 6 *Decision.* Choose the signal condition with the largest held-out $\hat{\Delta}_R(j)$ subject to predeclared gates; use $\mathcal{J}$, frontier scores, and allocation regret as diagnostics, not as the final deployment objective.

227 4.1 Progressive signals and paired router value

228 Let $Z_j(X; \zeta_j)$ be the router-visible pre-deployment record under signal condition j , where ζ_j denotes
 229 measurement randomness. Let Z_0 be the static signal observed before choosing a deployment action.
 230 For any admissible richer condition j , the router observes Z_j , pays the corresponding measurement
 231 and router-context cost, and selects an action

$$a_R^j(X) = R(Z_j(X; \zeta_j); U_R) \in \mathcal{A}.$$

232 Here U_R denotes router-side randomness; for deterministic routers it is fixed and suppressed. Thus
 233 the experimental comparison is between routing policies ρ_0, ρ_j induced by different signals for the
 234 same orchestrator R , not between different orchestrators. The normalized measurement loss is

$$\ell_{\text{meas}}(j, x; \zeta_j) = \lambda_{\text{meas}}^\top d_j(x; \zeta_j),$$

235 where d_j contains measurement wall-clock, tokens, checker calls, context expansion, parsing, retry,
 236 and latency costs. Set $\ell_{\text{meas}}(0, x) = 0$. The weights λ_{meas} use the same predeclared normalizers
 237 as deployment resource cost unless a separate router-cost scale is declared. In the main protocol,
 238 pre-deployment probes are outside the selected action's deployment run: they are charged through
 239 ℓ_{meas} and do not reduce the later horizon H or verifier budget B_V .

240 The selected-action deployment objective is

$$\mathcal{J}_R(j) := \mathbb{E} \left[\ell_{\text{meas}}(j, X; \zeta_j) + \ell_{\text{dep}}(a_R^j(X), X; \xi) \right]. \quad (7)$$

241 The expectation averages over task instances, measurement randomness, router randomness, and
 242 deployment randomness unless a conditional estimand is explicitly declared. The router-facing value
 243 of signal condition j is the paired gain

$$\Delta_R(j) := \mathcal{J}_R(0) - \mathcal{J}_R(j) = \mathbb{E} \left[\ell_{\text{dep}}(a_R^0(X), X) - \ell_{\text{dep}}(a_R^j(X), X) - \ell_{\text{meas}}(j, X) \right]. \quad (8)$$

244 Positive $\Delta_R(j)$ means that the richer signal improves net deployment loss for the same router on the
 245 same task distribution. Allocation shift alone is insufficient:

$$\text{Shift}_R(j) = \Pr[a_R^j(X) \neq a_R^0(X)]$$

246 measures decision relevance, not benefit.

247 **Proposition 1** (Unbiased paired router-gain estimator). *Assume task instances are sampled i.i.d. from*
 248 *μ , the orchestrator, action menu, and signal protocols are fixed before evaluation, measurement costs*
 249 *are recorded without bias, and deployment seeds are independent conditional on selected action and*
 250 *instance. The router is run once per instance and signal condition; the Q_{dep} deployment repeats*
 251 *estimate the conditional expected loss of the selected action. For condition j , estimate*

$$\hat{\Delta}_R(j) = \frac{1}{N} \sum_{i=1}^N \left[\frac{1}{Q_{\text{dep}}} \sum_{q=1}^{Q_{\text{dep}}} \ell_{\text{dep}}(a_{R,i}^0, x_i; \xi_{i,q}^0) - \frac{1}{Q_{\text{dep}}} \sum_{q=1}^{Q_{\text{dep}}} \ell_{\text{dep}}(a_{R,i}^j, x_i; \xi_{i,q}^j) - \hat{\ell}_{\text{meas}}(j, x_i) \right].$$

252 *If the repeated-run averages are unbiased for the expected deployment losses of the selected actions,*
 253 *then $\mathbb{E}[\hat{\Delta}_R(j)] = \Delta_R(j)$.*

254 When the same worker is selected under Z_0 and Z_j , using common deployment seeds preserves
 255 unbiasedness and reduces paired variance. When counterfactual outcomes for all actions are logged
 256 on the same instance, hindsight action regret can be reported as a diagnostic, but it is not the primary
 257 estimand. If the study fixes a finite prototype panel rather than sampling i.i.d. instances from μ , the
 258 same estimator describes that panel and should not be interpreted as a population-level deployment
 259 estimand.

260 5 Operational experimental protocol

261 This section specifies the controlled protocol used to test the routed-policy decomposition. The
 262 main study is a fixed-prototype repeated-run panel: it fixes three AutoResearch task modes and uses
 263 repeated agentic trajectories to estimate, for the same starting executable artifact, how worker choice,
 264 signal exposure, and routing affect first-passage success and cost. Accordingly, the primary estimates
 265 are panel analogues of the population quantities in Section 3. Aggregate panel summaries use the
 266 predeclared empirical mode weighting $\hat{\pi}(s) = 1/3$ for $s \in \mathcal{S}$.

267 **Task modes.** The task is AutoResearch-style executable-artifact optimization on CIFAR-10. Each
 268 task instance contains an editable `train.py`, a read-only data/checker module `prepare.py`, and
 269 a verifier that executes `train.py` and reports task-level diagnostics for validation performance,
 270 compute cost, optimization progress, and model complexity. In this protocol, each task mode is
 271 represented by one fixed executable prototype of the same CIFAR-10 task. Concretely, a task mode is
 272 specified by the tuple

$$\chi_s = (A_s, D, \text{split}, u, \eta_s, V, B_V, H, \delta),$$

273 where A_s is the starting model class and editable `train.py` template, D is CIFAR-10, `split` is the
 274 fixed train/validation split, u is the model-initialization seed, η_s denotes starting optimization and
 275 architecture hyperparameters, V is the checker, B_V is the checker budget, H is the agent horizon,
 276 and δ is the success threshold. The protocol uses three task modes:

$$\mathcal{S} = \{\text{mlp_flat}, \text{cnn_compact}, \text{resnet_micro}\}.$$

277 **Execution protocol.** Appendix D reports the full fixed configuration read from the current
 278 `train.py` templates and the CIFAR-10 checker in Table 5. The paper-facing worker menu is

$$\mathcal{M}_{\text{main}} = \{\text{gpt_5_4_mini}, \text{gpt_5_3_codex}, \text{gpt_5_4}\}.$$

279 In the main experiment, each deployable action is one worker under the common fixed prompt
 280 and harness, so $\mathcal{A} = \mathcal{M}$. In the current snapshot, `gpt_5_3_codex` and `gpt_5_4` are the mature
 281 comparison pair. `gpt_5_4_mini` is being filled in as the lower-cost replacement for the unavailable
 282 `gpt_5_3_codex_spark` condition. Each worker uses the same prompt. For every task-mode-worker
 283 cell, we run

$$n_{\text{cell}} = 10$$

284 independent trajectories from the same fixed `train.py`; the repeated-run index changes only the
 285 agentic trajectory. The router is a fixed prompt-only orchestrator that chooses one worker from $\mathcal{M}$
 286 before deployment. For the primary protocol it is deterministic after fixing model identifier, prompt,
 287 parser, temperature, and seed, so $B_R = 1$ router draw is logged per task-mode-signal pair. We
 288 evaluate a sequential signal protocol

$$Z_0, Z_1, Z_2, Z_3.$$

289 The signals are nested: Z_j contains all fields of Z_{j-1} plus the additional fields declared for condition
 290 j . Z_0 contains only the task name, checker objective, worker menu, and budget. Z_1 adds the initial
 291 `train.py` and a static workload summary: starting model class, parameter count, optimizer and
 292 budget settings, dataset/checker identifiers, and parse/runtime availability. Z_2 adds a low-cost checker
 293 probe of the unmodified artifact at a predeclared probe budget. Z_3 adds the full unmodified-artifact
 294 baseline under the deployment checker budget. Probe records contain validation loss, validation
 295 accuracy, training seconds, total seconds, optimizer steps, parameter count, parse/runtime status,
 296 and the probe budget. Probe runs are measurement actions for the router only: they are charged
 297 through ℓ_{meas} and do not consume the selected worker’s later deployment horizon or checker budget.
 298 Let $L_{\text{base}}(x)$ be the validation loss of the unmodified starting artifact for instance x under the fixed
 299 deployment checker. This evaluator-only baseline is used to score G_t in all conditions; it becomes
 300 router-visible only in Z_3 and is then charged as measurement information. In the fixed-prototype
 301 protocol this baseline is constant within each task mode. At edit step t , a worker produces a candidate
 302 `train.py`, the checker returns loss L_t , and checked progress is

$$G_t = \text{clip}_{[0,1]} \left(\frac{L_{\text{base}}(x) - L_t}{|L_{\text{base}}(x)| + \varepsilon_L} \right).$$

303 We set $\varepsilon_L = 10^{-8}$ in the primary analysis. The primary success event is first passage above the
 304 relative-improvement threshold

$$\tau_\delta = \inf\{t \leq H : G_t \geq \delta\}, \quad \delta = 0.05.$$

305 The fixed-prototype competence estimator is $\hat{p}_{\chi_s}(a) = n_{\text{cell}}^{-1} \sum_{r=1}^{n_{\text{cell}}} \mathbf{1}\{\tau_{\delta,a,s,r} \leq H\}$, the panel
 306 analogue of $p(a, s)$. The primary cost $\hat{\kappa}_{\chi_s}(a)$ is the sample mean worker wall-clock accumulated to
 307 $\tau_\delta \wedge H$, the panel analogue of $\kappa(a, s)$; token count is reported as a sensitivity analysis. The primary
 308 deployment loss is the first-passage loss in (5); threshold occupancy and final checked quality are
 309 reported through the audit loss (6). For log scores, we use the predeclared smoothed success estimate
 310 $\tilde{p}_{\chi_s}(a)$ to avoid infinite values in zero-success pilot cells.

311 **Experiment 1: worker frontier.** Run all workers on all three task modes without a router:

$$3 \text{ task modes} \times 3 \text{ workers} \times n_{\text{cell}} = 90$$

312 full trajectories. Estimate $\hat{p}_{\chi_s}(a)$, $\hat{\kappa}_{\chi_s}(a)$, mean first-hit step, occupancy, final quality, and uncertainty
 313 intervals. For each task mode, construct the empirical cost-success frontier

$$\hat{\mathcal{F}}(s) = \left\{ a \in \mathcal{A} : \log \hat{\kappa}_{\chi_s}(a) - \log \tilde{p}_{\chi_s}(a) \leq \min_{a' \in \mathcal{A}} [\log \hat{\kappa}_{\chi_s}(a') - \log \tilde{p}_{\chi_s}(a')] + \epsilon_{\text{frontier}} \right\},$$

314 where $\epsilon_{\text{frontier}}$ is a predeclared uncertainty tolerance. This experiment is the worker-side calibration
 315 for Theorem 1. It estimates fixed-prototype analogues of $\kappa(a, s)$ and $p(a, s)$ for every candidate action
 316 before any router is allowed to choose. The same estimates can be plugged into Theorem 2, because
 317 the first-passage success probabilities and first-passage costs determine when repeated deployments
 318 of a lower-cost worker can match the one-attempt success probability of a stronger worker.




Figure 1: Mature-worker snapshot on common pilot support, excluding the rate-limited Spark condition. Left: deployment loss, where lower is better and intervals are the bootstrap intervals produced by the accounting script. Right: the certified log-effort surrogate used by the decomposition. The key qualitative result is a rank disagreement: gpt_5_4 reaches the threshold faster and achieves higher final improvement, but the deployment-loss accounting currently favors gpt_5_3_codex in all three modes, and the log-effort surrogate favors different workers by mode.

Table 1: Pilot worker-frontier summary on common support. n_{common} is the number of paired trial indices retained for the two mature workers in that mode. The table reports first-hit success at $\delta = 0.05$, mean first-hit step among successful runs, mean threshold occupancy, mean final relative improvement, the log-effort surrogate, and deployment loss. Lower is better for the last two columns.

Mode	Worker	n_{common}	Success	Mean τ	Occupancy	Final improv.	Log-effort	Dep. loss
cnn_compact	gpt_5_3_codex	10	10/10	6.50	0.725	0.232	5.814	0.848
cnn_compact	gpt_5_4	10	10/10	3.30	0.885	0.373	5.503	1.580
mlp_flat	gpt_5_3_codex	10	9/10	8.11	0.580	0.171	6.045	0.890
mlp_flat	gpt_5_4	10	10/10	2.60	0.920	0.267	5.358	1.158
resnet_micro	gpt_5_3_codex	9	10/10	3.20	0.890	0.194	5.102	0.962
resnet_micro	gpt_5_4	9	9/9	1.67	0.967	0.282	5.626	1.934

Current worker-frontier results. Because the Spark condition hit rate limits, the main snapshot below compares the two mature workers only: gpt_5_3_codex and gpt_5_4. For each mode, we use the maximum common support available to both workers: ten pilot trajectories for cnn_compact, ten for mlp_flat, and nine for resnet_micro. The provisional gpt_5_4_mini pilot is intentionally not pooled into these estimates; it is reserved for the final lower-cost comparison once the cell reaches the same support.

This is the first empirical support for the paper’s main claim. If one ranks workers by raw final improvement or by first-hit speed, gpt_5_4 looks dominant. If one ranks by the deployment loss currently charged by the protocol, the recommendation changes. The experiment therefore already exhibits the failure mode of scalar leaderboard racing: different components of time-to-certified-solution point to different engineering choices.

Experiment 2: router-response audit. For each task mode and each signal Z_j , query the same orchestrator under the same worker menu. The router returns a structured decision record containing selected worker, confidence, short task diagnosis, evidence used, expected failure mode, and expected cost risk. The diagnostic question is whether progressive signals change the router’s own decision record: selected-worker distribution, confidence, diagnosis consistency with the declared task mode, and cited evidence. This is a router-response diagnostic, not an estimator of $I_\pi(S; Z)$. An information estimator would require a predeclared finite representation or predictive model for the signal records.

Current signal-audit results. The current router-facing diagnostic is not yet a deployment claim. It asks the narrower question required by the information term: whether admissible pre-deployment




Figure 2: Threshold sensitivity for the same mature-worker comparison. The left panel shows first-passage success probability as the required relative-improvement threshold varies; the right panel shows the median first-hit step. At the paper’s operating point $\delta = 0.05$, gpt_5_4 has higher success probability and lower median first-hit time than gpt_5_3_codex. The fact that this speed advantage does not automatically imply lower deployment loss is the operational distinction the decomposition is meant to expose.




Figure 3: Applicability check for the retry-crossover theorem on the frozen two-worker confirmation split. The theorem requires a lower-cost action with lower one-attempt success probability and a stronger but more expensive action. The two mature workers do not cleanly enter this regime: gpt_5_4 has both higher success and lower first-passage cost on mlp_flat, while both workers succeed on every frozen resnet_micro run and gpt_5_3_codex is cheaper there. The crossover result is therefore a design rule for adding a genuinely lower-cost worker, not a promoted empirical claim of the present two-worker menu.

339 records distinguish the task modes that have different empirical frontiers. On the current snapshot,
 340 budget-only records collapse the three modes to the same predicted class, while probe-containing
 341 records perfectly recover the predeclared mode labels.

342 This supports the premise that low-cost probes can reveal mode information, but it does not yet
 343 prove that the router uses that information correctly. The deployment-facing question is deferred to
 344 Experiments 3 and 4: whether the induced worker choice moves toward the empirical frontier and
 345 pays for the measurement cost.

346 **Experiment 3: routing shift and allocation regret.** Using the decisions from Experiment 2,
 347 compute the empirical panel shift rate

$$\widehat{\text{Shift}}_R(j) = \sum_{s \in \mathcal{S}} \hat{\pi}(s) \mathbf{1}\{\rho_j(\chi_s) \neq \rho_0(\chi_s)\}$$

Table 2: Router-facing mode-diagnosis audit on the current signal records. Accuracy is reported only as a diagnostic for whether the signal exposes the finite task-mode structure; it is not the deployment objective.

Signal record	Records	Accuracy	Macro accuracy
Budget only	63	0.349	0.333
Probe only	63	1.000	1.000
Probe + budget	63	1.000	1.000
Leaky current-state control	63	1.000	1.000




Figure 4: Final two-worker router-response audit. Left: selected-worker distribution under real signal records. Right: allocation regret against the frozen log-effort frontier. Richer records change the router only weakly and do not monotonically reduce regret, even though the signal audit shows that probes contain mode information.

for the deterministic primary router. Then use the frontier from Experiment 1 to score each selected action by allocation regret

$$\hat{r}_j(\chi_s) = [\log \hat{\kappa}_{\chi_s}(\rho_j(\chi_s)) - \log \tilde{p}_{\chi_s}(\rho_j(\chi_s))] - \min_{a \in \mathcal{A}} [\log \hat{\kappa}_{\chi_s}(a) - \log \tilde{p}_{\chi_s}(a)].$$

This is an operational action-frontier diagnostic, not an estimator of the KL mode-summary divergence $\mathbb{E}_Z \text{KL}(\pi_Z \| q_Z)$ unless a separate construction of q_Z has been predeclared. When the same pilot runs define the frontier and score regret, regret is reported as an in-sample diagnostic; cross-fitted regret is used whenever enough pilot trajectories are available. Shift rate measures whether Z_j changes the routing policy, while allocation regret asks whether the changed policy moves toward the empirical mode-conditional frontier from Experiment 1. A signal can induce a router response and still increase allocation regret if it makes the orchestrator change worker in the wrong direction. Conversely, low regret indicates that the selected action moved closer to the empirical mode-conditional frontier rather than merely increasing decision variability.

Current routing-regret status. The final two-worker router audit contains 480 prompt-only decisions: 30 instances, four signal levels, and four control conditions. On the real signal records, the router remains strongly biased toward `gpt_5_4`: it selects `gpt_5_4` on 27/30 budget-only records, 29/30 static-workload records, 30/30 probe records, and 28/30 full-scout records. Mean log-effort regret does not decrease with richer information: it is 0.166 for Z_0 , 0.194 for Z_1 , 0.158 for Z_2 , and 0.197 for Z_3 . Thus the router-response audit separates information availability from useful allocation. The probe signals expose the predeclared task modes, but the prompt-only router mostly maps them to the same worker and does not reliably move toward the cost-adjusted frontier in Figure 5.

Experiment 4: paired selected-worker deployment gain. Report all $j \in \{1, 2, 3\}$ with multiplicity control as the primary analysis. If a single non-static signal is selected on a diagnostic split, evaluate it only on a separate holdout split. For each task mode and paired deployment index r , compare the worker selected under Z_0 with the worker selected under Z_j from the same original artifact and horizon. Pairing is by common prototype, common checker configuration, and matched deployment seed when the backend exposes seed control; otherwise intervals are stratified by task




Figure 5: Frozen two-worker confirmation frontier with 25 holdout trajectories per task-mode-worker cell. The left panel reports deployment loss with bootstrap intervals; the right panel reports the certified log-effort surrogate. The confirmation split preserves the qualitative rank disagreement from the pilot: gpt_5_4 has lower first-hit time, but deployment loss favors gpt_5_3_codex on cnn_compact and resnet_micro, while mlp_flat is nearly tied.

Table 3: Frozen confirmation summary for the two mature workers. Each cell uses the first 25 materialized holdout seeds under the frozen deterministic seed order. Deployment-loss intervals are bootstrap intervals over the 25 trajectories.

Mode	Worker	Success	Mean τ	Dep. loss	95% CI	Log-effort
cnn_compact	gpt_5_3_codex	24/25	5.83	1.011	[0.843, 1.231]	5.577
cnn_compact	gpt_5_4	25/25	2.96	1.438	[1.170, 1.770]	5.498
mlp_flat	gpt_5_3_codex	23/25	7.13	1.065	[0.867, 1.283]	6.028
mlp_flat	gpt_5_4	25/25	2.44	1.053	[0.901, 1.238]	4.955
resnet_micro	gpt_5_3_codex	25/25	2.36	0.997	[0.894, 1.105]	4.981
resnet_micro	gpt_5_4	25/25	1.84	2.193	[1.593, 2.996]	5.456

mode rather than treated as seed-paired. The selected-action confirmation uses $Q_{\text{dep}} = 10$ repeats per selected action and prototype unless a larger holdout budget is declared. Estimate

$$\hat{\Delta}_R(j) = \hat{\ell}_{\text{dep}}(\rho_0) - \hat{\ell}_{\text{dep}}(\rho_j) - \hat{\ell}_{\text{meas}}(j).$$

Report confidence intervals over the declared pairing unit, first-pass success changes, cost-to-success changes, gross loss reduction, measurement cost, and the subset where the signal changed the worker but increased loss. This experiment is the deployment validation of the theory-facing diagnostics. Proposition 1 gives the estimand and estimator for the same-orchestrator comparison, while (5) defines the loss used to decide whether a signal-conditioned policy is actually better after measurement cost is charged. Positive routing information and lower allocation regret are not sufficient by themselves; the routed-policy claim is supported only when the paired gain $\hat{\Delta}_R(j)$ is positive on held-out trajectories.

Current validation status. For the two mature workers, we freeze the confirmation analysis at 25 runs per task-mode-worker cell:

$$3 \text{ task modes} \times 2 \text{ workers} \times 25 = 150$$

holdout trajectories. The frozen comparison uses gpt_5_3_codex and gpt_5_4; Spark is excluded because the endpoint hit rate limits, and gpt_5_4_mini is reserved for a later lower-cost extension once its cells reach comparable support. The paired-gain analysis therefore evaluates whether router-visible information improves deployment loss within this two-worker action menu.

Paired-gain result. The paired deployment validation is negative for the current prompt-only router. Against Z_0 , richer real signals produce small shift rates and negative net gain after measurement cost: Z_1 shifts 13.3% of decisions with net gain -0.043 and bootstrap interval $[-0.100, -0.000]$; Z_2 shifts 10.0% with net gain -0.050 , interval $[-0.107, -0.007]$; and Z_3 shifts 10.0% with net gain -0.057 ,




Figure 6: Paired deployment gain for real signal records relative to Z_0 , using the frozen two-worker holdout frontier. All richer signal levels have negative net gain after measurement cost, so the current prompt-only router fails the deployment claim gate.

Table 4: Paired router validation on the final two-worker menu. Net gain is $\hat{\ell}_{\text{dep}}(\rho_0) - \hat{\ell}_{\text{dep}}(\rho_j) - \hat{\ell}_{\text{meas}}(j)$; positive is better.

Signal	Pairs	Shift	Gross gain	Meas. loss	Net gain
Z_1	30	0.133	-0.043	0.000	-0.043
Z_2	30	0.100	-0.043	0.007	-0.050
Z_3	30	0.100	-0.029	0.028	-0.057

interval $[-0.100, -0.028]$. The gross gains are already non-positive, and measurement cost makes the probe/scout conditions worse. This is the critical falsification result for the deployment-facing claim: informative signals are not sufficient; the deployed router must use them in a way that improves the selected action’s checked cost–success tradeoff.

Experiment 5: negative controls. Repeat Experiments 2–4 with schema-preserving controls: shuffled probe records, probes from the wrong task mode, and synthetic noise records with the same field names and token length. The same router, prompt, worker menu, and scoring rule are used. These controls test whether router response, allocation regret, and paired deployment gain are specific to task-relevant signal content. They preserve the format, approximate context burden, and measurement/context cost of the corresponding real Z_j while breaking its relationship to the task mode and worker frontier. A credible deployment claim requires the real signal to reduce allocation regret and improve paired deployment gain more than these controls after measurement cost is charged.

Current negative-control status. The negative controls are also complete on the final two-worker menu. They do not show a clean separation between real task-relevant signals and corrupted signals. For example, wrong-mode Z_2 has a small positive mean net gain (0.017) with an interval crossing zero, while several shuffled or synthetic controls are negative with intervals similar to the real signals. The right interpretation is not that the controls discover a useful policy; it is that the prompt-only router has weak, noisy response to the records and the real signal does not dominate the controls. Accordingly, Experiment 5 blocks promotion of a strong routing-information claim. The paper should promote the worker-frontier and decomposition diagnostics, and report the router experiment as a falsification-style result for this particular router.

Interpretation through the two theorems. The four-term identity is empirically useful as a diagnostic accounting language. The worker-side terms show nontrivial cost and competence structure by mode; the signal audit shows that Z can expose task-mode information; the router audit then shows that information value is lost through allocation behavior. In the theorem’s terms, positive routing information is not enough: allocation mismatch and worker-side costs can erase it, and the final deployment loss confirms that erasure. The retry-crossover theorem gives a complementary design rule, but the two mature workers do not instantiate its assumptions. There is no stable “cheap




Figure 7: Negative-control paired gains for shuffled probes, wrong-mode probes, and synthetic noise records. The real records do not cleanly outperform the controls, which indicates that the current prompt-only router is not reliably exploiting task-relevant measurement content.

but weaker” action in the final two-worker menu: gpt_5_4 is both stronger and cheaper to first passage on mlp_flat, while gpt_5_3_codex is cheaper on resnet_micro with no success-rate penalty. This is why gpt_5_4_mini is scientifically useful as a later extension: it would test the retry-crossover regime directly rather than merely adding another strong-worker comparison.

Claim promotion gates. We promote a routing recommendation only when the following are reported: mode-conditional competence with uncertainty, worker-side cost sensitivity, router shift, allocation regret, paired deployment gain with measurement cost charged, negative controls, calibration of $\tilde{p}$, and surrogate-versus-loss ranking agreement. If routing diagnostics improve but paired deployment gain does not, the conclusion is diagnostic rather than prescriptive. If paired deployment gain improves but surrogate-versus-loss agreement is poor, the protocol found a useful router but did not validate the packed-mode explanation.

6 Related work

The closest classical ancestor is algorithm selection. Since Rice’s formulation, solver portfolios have treated performance as instance-dependent: SATzilla and AutoFolio learn to select solvers from instance features rather than report one globally best solver [5–8]. That literature also includes feature-computation costs, pre-solving, and probing. Our contribution is the specialization to externally checked LLM agents, where a router observes a controlled pre-deployment record, selects a worker or harness action, and is scored by paired deployment loss rather than by standalone routing accuracy. This differs from AutoML systems such as auto-sklearn, which select and tune learning pipelines offline for a dataset; here the unit is a deployable action for a concrete verifier-backed instance [9].

LLM routing and cascades provide the model-menu context. FrugalGPT frames model cascades as cost reduction across heterogeneous APIs, RouteLLM learns cost-quality routing from preference data, and recent compound-system routing work extends the idea to multi-module systems [10–14]. RouterBench and LLMRouterBench make this evaluation problem explicit at benchmark scale, with cost-, latency-, and quality-aware routing outcomes over large multi-model pools [15, 16]. These methods establish that heterogeneous model menus can create economic gains. The present paper asks how much checked, task-specific information is worth before committing to an expensive agentic deployment.

Checked agent benchmarks provide the task substrate. SWE-bench evaluates repository repair against tests, AutoResearch evaluates training-loop edits against validation feedback, and theorem-proving systems use analogous external validators [1, 17]. Most benchmarks ask how well a worker performs after it is launched. Our question is upstream: which worker should be launched, what should the router be allowed to observe, and how should the cost of that observation be charged?

454 Finally, repeated sampling and inference-time scaling show that lower-cost or smaller models can
 455 become competitive when verification is available [18–21]. The first-passage retry crossover in
 456 Theorem 2 is the simplest version of that tradeoff. The four-term identity embeds the same idea in a
 457 mode-aware routing and measurement account.

458 7 Discussion, future directions, and conclusion

459 There are two levels of claim. The theory-facing claim is an accounting identity under Assumption 1.
 460 It is exact, but only for the certified log-effort surrogate. The deployment-facing claim is empirical:
 461 the accounting should predict or explain useful action selection under a richer deployment loss. These
 462 two claims should not be collapsed. The current experiments support the central diagnostic claim:
 463 raw worker rankings, first-passage rankings, log-effort rankings, and deployment-loss rankings can
 464 disagree on the same checked task family. That disagreement is precisely the regime in which a
 465 budget-aware orchestration framework is needed. They do not support the stronger claim that the
 466 current prompt-only router improves deployment loss. Instead, the router validation is a useful
 467 negative result: task-relevant information is measurable, but a router that mostly defaults to the
 468 stronger model can still lose after measurement cost and worker-side costs are charged.

469 The packed-mode assumption is the main modeling limitation. Real agent trajectories are stateful,
 470 attempts are not independent, and task modes may be fuzzy. The paper therefore reports first-passage
 471 curves, occupancy, hazard diagnostics, cost sensitivity, and surrogate-versus-loss ranking agreement
 472 instead of assuming that the finite-mode surrogate is correct. If these diagnostics are unstable, the
 473 decomposition remains a useful diagnostic language but not a reliable selection rule.

474 The operational task-mode schema is also a limitation. If S is too coarse, mode-conditional compe-
 475 tence is washed out and the router appears useless. If S is too fine, sample sizes per cell collapse and
 476 the signal-mode estimates become unstable. The present evaluation fixes the executable prototypes
 477 and mode schema to isolate agentic trajectory variance; task-instance and schema perturbations are
 478 left as robustness extensions rather than promoted experiments.

479 Finally, measurement value is router-specific. A baseline probe can contain information about the
 480 best action and still fail to improve deployment loss if the actual router does not use it well or if
 481 the measurement is too expensive. For this reason, the paired gain $\Delta_R(j)$ is the main deployment
 482 estimand, while mutual-information diagnostics, mode-summary divergence, and allocation regret
 483 help explain why that gain appears or fails to appear. This distinction is the main empirical lesson of
 484 the final two-worker study. The Z_2 and Z_3 records contain task-mode information, but the induced
 485 policies have negative paired gain. The decomposition therefore functions as a claim discipline: it
 486 prevents us from promoting information availability as deployment value.

487 **Future directions.** The immediate experimental extension is to finish the `gpt_5_4_mini` pilot and
 488 holdout cells, replacing the unavailable Spark worker with a stable lower-cost action. That would
 489 make the worker menu span a broader cost–capability frontier than the two-worker mature comparison
 490 reported here. It would also test the retry-crossover theorem in its intended regime, where a lower-cost
 491 worker may compensate for lower one-attempt success through repeated checked attempts. Beyond
 492 this benchmark, the same protocol should be tested on repository repair, theorem checking, and
 493 multi-file research-code editing, where the mode schema is less artificial and measurement costs are
 494 more heterogeneous.

495 **Conclusion.** Model choice inside an agent harness is not a scalar leaderboard problem. It is a
 496 measurable decomposition problem validated by paired deployment outcomes. The right worker
 497 depends on task state, checked competence, routing information, allocation quality, first-passage time,
 498 retry depth, and worker-side cost. Our CIFAR-10 edit–verify campaign provides the first empirical
 499 slice of that argument: the same two mature workers can trade off speed, final improvement, log-effort,
 500 and deployment loss in different ways. The two-worker routing analysis asks the deployment-facing
 501 question directly and answers it negatively for the current prompt-only router: pre-deployment
 502 information does not move allocation reliably enough to improve paired deployment loss after its
 503 measurement cost is charged. That negative result is not a failure of the framework; it is the framework
 504 doing its job, distinguishing worker-frontier structure from a validated routing recommendation.

References

- [1] Carlos E. Jimenez, John Yang, Alexander Wettig, Shunyu Yao, Kexin Pei, Ofir Press, and Karthik Narasimhan. SWE-bench: Can language models resolve real-world GitHub issues? In *International Conference on Learning Representations (ICLR)*, 2024.
- [2] Jun Shern Chan, Neil Chowdhury, Oliver Jaffe, James Aung, Dane Sherburn, Evan Mays, Giulio Starace, Kevin Liu, Leon Maksin, Tejal Patwardhan, Lilian Weng, and Aleksander Madry. MLE-bench: Evaluating machine learning agents on machine learning engineering. *arXiv preprint arXiv:2410.07095*, 2024.
- [3] Mahmoud Abdelmoneum, Pierfrancesco Beneventano, and Tomaso Poggio. PoggioAI/MSc: ML theory research with humans on the loop. Technical Report Technical Report v0, MIT, 2026.
- [4] Alessandro Achille and Stefano Soatto. AI agents as universal task solvers: It’s all about time. *Entropy*, 28(3):332, 2026.
- [5] John R. Rice. The algorithm selection problem. In *Advances in Computers*, volume 15, pages 65–118. Elsevier, 1976.
- [6] Lin Xu, Frank Hutter, Holger H. Hoos, and Kevin Leyton-Brown. SATzilla: Portfolio-based algorithm selection for SAT. *Journal of Artificial Intelligence Research*, 32:565–606, 2008.
- [7] Marius Lindauer, Katharina Eggersperger, Matthias Feurer, Bernd Bischl, and Frank Hutter. AutoFolio: An automatically configured algorithm selector. *Journal of Artificial Intelligence Research*, 53:745–778, 2015.
- [8] Lars Kotthoff. Algorithm selection for combinatorial search problems: A survey. *AI Magazine*, 35(3):48–60, 2016.
- [9] Matthias Feurer, Aaron Klein, Katharina Eggersperger, Jost Tobias Springenberg, Manuel Blum, and Frank Hutter. Efficient and robust automated machine learning. In *Advances in Neural Information Processing Systems*, volume 28, 2015.
- [10] Lingjiao Chen, Matei Zaharia, and James Zou. FrugalGPT: How to use large language models while reducing cost and improving performance. *arXiv preprint arXiv:2305.05176*, 2023.
- [11] Isaac Ong, Amjad Almahairi, Vincent Wu, Wei-Lin Chiang, Tianhao Wu, Joseph E. Gonzalez, M Waleed Kadous, and Ion Stoica. RouteLLM: Learning to route LLMs with preference data. *arXiv preprint arXiv:2406.18665*, 2024.
- [12] Lingjiao Chen et al. Optimizing model selection for compound AI systems. *arXiv preprint arXiv:2502.14815*, 2025.
- [13] Yanwei Yue, Guibin Zhang, Boyang Liu, Guancheng Wan, Kun Wang, Dawei Cheng, and Yiyang Qi. MasRouter: Learning to route LLMs for multi-agent systems. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15549–15572, 2025.
- [14] Matei Zaharia, Omar Khattab, Lingjiao Chen, Jared Quincy Davis, Heather Miller, Chris Potts, James Zou, Michael Carbin, Jonathan Frankle, Naveen Rao, and Ali Ghodsi. The shift from models to compound AI systems. Berkeley Artificial Intelligence Research Blog, 2024.
- [15] Qitian Jason Hu, Jacob Bieker, Xiuyu Li, Nan Jiang, Benjamin Keigwin, Gaurav Ranganath, Kurt Keutzer, and Shriyash Kaustubh Upadhyay. RouterBench: A benchmark for multi-LLM routing system. *arXiv preprint arXiv:2403.12031*, 2024.
- [16] Anonymous. LLMRouterBench: A massive benchmark and unified framework for LLM routing, 2026. Anonymous ACL submission.
- [17] Andrej Karpathy. autoresearch: AI agents running research on single-GPU nanochat training automatically. <https://github.com/karpathy/autoresearch>, 2026. GitHub repository.

- [18] Bradley Brown, Jordan Juravsky, Ryan Ehrlich, Ronald Clark, Quoc V. Le, Christopher Ré, and Azalia Mirhoseini. Large language monkeys: Scaling inference compute with repeated sampling. *arXiv preprint arXiv:2407.21787*, 2024.
- [19] Charlie Victor Snell, Jaehoon Lee, Kelvin Xu, and Aviral Kumar. Scaling LLM test-time compute optimally can be more effective than scaling model parameters. In *International Conference on Learning Representations (ICLR)*, 2025.
- [20] Yangzhen Wu, Zhiqing Sun, Shanda Li, Sean Welleck, and Yiming Yang. Inference scaling laws: An empirical analysis of compute-optimal inference for problem-solving with language models. In *International Conference on Learning Representations (ICLR)*, 2025.
- [21] Jinghan Zhang, Kunhao Zheng, Yici Yan, Yuchen Ouyang, and Jun Zhu. Scaling LLM inference with optimized sample compute allocation. *arXiv preprint arXiv:2410.22480*, 2024.

A Proof details

A.1 Four-term identity

Under Assumption 1,

$$\log T_{\rho,q}(S, Z) = C + \log \kappa(a_Z, S) - \log p(a_Z, S) - \log q_Z(S).$$

Conditioning on $Z = z$,

$$\mathbb{E}[-\log q_z(S) \mid Z = z] = \sum_s \pi_z(s) \log \frac{1}{q_z(s)} = H(\pi_z) + \text{KL}(\pi_z \parallel q_z).$$

Therefore

$$\mathbb{E}[\log T_{\rho,q}] = C + \mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} [\log \kappa(a_Z, S) - \log p(a_Z, S)] + \mathbb{E}_Z [H(\pi_Z)] + \mathbb{E}_Z [\text{KL}(\pi_Z \parallel q_Z)].$$

For the prior-matched baseline ρ_0 that deploys fixed action a_0 and uses $q_0(\cdot \mid z) \equiv \pi$,

$$\mathbb{E}[\log T_{\rho_0, \pi}] = C + \mathbb{E}_{S \sim \pi} [\log \kappa(a_0, S) - \log p(a_0, S)] + H(\pi).$$

Subtracting the deployed effort from the baseline effort yields

$$\Delta = \mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} \left[\log \frac{\kappa(a_0, S)}{\kappa(a_Z, S)} \right] + \mathbb{E}_Z \mathbb{E}_{S \sim \pi_Z} \left[\log \frac{p(a_Z, S)}{p(a_0, S)} \right] + H(\pi) - \mathbb{E}_Z [H(\pi_Z)] - \mathbb{E}_Z [\text{KL}(\pi_Z \parallel q_Z)].$$

The entropy difference is $I_\pi(S; Z)$, giving (4).

A.2 Pilot concentration

For binary success observations in an action-mode cell, Hoeffding gives

$$\mathbb{P}(|\hat{p}(a, s) - p(a, s)| > \alpha) \leq 2 \exp(-2n_0 \alpha^2).$$

If all true probabilities are at least $p_{\min}$, then the log-probability error obeys $|\log \hat{p} - \log p| \leq \alpha/p_{\min} + O(\alpha^2/p_{\min}^2)$ on the good event. A union bound over $|\mathcal{A}| |\mathcal{S}|$ cells yields the conservative pilot rule

$$n_0 = O\left(\frac{\log(|\mathcal{A}| |\mathcal{S}| / \delta_{\text{conf}})}{\delta_{\text{acc}}^2 p_{\min}^2}\right).$$

Wilson intervals and bootstrap intervals are reported in practice because pilot cell counts are small.

B Operational diagnostics for the four terms

The identity above becomes useful only after each term is tied to quantities logged by the deployment protocol. The checker-facing progress process is the common object used to operationalize competence, cost, information, and allocation diagnostics. Fix a horizon H and a task-normalized checked progress process $G_t(a, x) \in [0, 1]$ for action a on instance x . In the main experiment, the router acts

once before deployment: after observing Z , it selects one action, and that action is used for the whole trajectory. The success and cost diagnostics below are therefore trajectory-level quantities computed from the step-wise verifier trace. For lower-is-better checker losses, one admissible normalization is

$$G_t(a, x) = \text{clip}_{[0,1]} \left(\frac{L_0(x) - L_t(a, x)}{|L_0(x)| + \varepsilon_L} \right),$$

where $L_0(x)$ is the starting-artifact loss and $L_t(a, x)$ is the checked loss convention at step t . Given a success threshold $\delta > 0$,

$$\tau_\delta(a, x) = \inf\{t \leq H : G_t(a, x) \geq \delta\}, \quad F_\delta(a, x) = \mathbf{1}\{\tau_\delta(a, x) = \infty\}.$$

We also log persistence through

$$\text{Occ}_{\delta,H}(a, x) = \frac{1}{H} \sum_{t=1}^H \mathbf{1}\{G_t(a, x) \geq \delta\}.$$

Competence. The action competence component $p(a, s)$ is the first-success probability of action a on task mode s :

$$p(a, s) = \mathbb{P}[\tau_\delta(a, X) \leq H \mid S = s].$$

Thus G_t is step-wise, but the policy competence term uses the probability that the worker selected by the router reaches the success threshold at least once before the horizon. Cumulative progress and persistence are not substituted for $p(a, s)$ in the identity. We log

$$\bar{G}_a(s) = \mathbb{E} \left[\frac{1}{H} \sum_{t=1}^H G_t(a, X) \mid S = s \right]$$

and $\text{Occ}_{\delta,H}$ as secondary diagnostics, especially when first-hit success saturates or when partial progress is scientifically relevant.

Cost. The action cost component $\kappa(a, s)$ is kept separate from competence. In the main protocol it is the normalized cost of deploying action a on mode s for the same trajectory convention used to define $p(a, s)$. The routed-policy cost term averages this mode-conditional cost over the actions selected under Z . The primary convention is cost accumulated until the first certified success or the horizon, $\tau_\delta \wedge H$. We log

$$C_\delta(a, x; \xi) = \lambda^\top (\tilde{T}_{\tau_\delta \wedge H}^{\text{worker}}, \tilde{C}_{\tau_\delta \wedge H}^{\text{tok}})(a, x; \xi),$$

and report worker wall-clock cost as primary, with token count as a sensitivity analysis. Checker runtime is logged for audit, but it is not part of $\kappa(a, s)$ unless the checker protocol or checker-call frequency differs across actions.

Routing information. The theorem term $I(S; Z)$ measures how much a router-visible signal Z separates the predeclared task modes before deployment. This does not mean that the evaluator already knows the best worker for each instance. The evaluator knows only the candidate mode schema and can estimate whether Z_j makes that schema more identifiable:

$$I(S; Z_j \mid Z_0) = H(S \mid Z_0) - H(S \mid Z_j).$$

The quantity is estimated only when a finite signal representation or calibrated predictive model for the signal records is predeclared. Otherwise the paper reports signal summaries and router-response diagnostics instead of a Shannon mutual-information estimate. It is a descriptive information diagnostic; its decision value is established only after action outcomes show that the modes predict different cost–success profiles. The paired router experiment then checks whether higher-information signals actually change the selected action and improve first-passage deployment outcomes.

Allocation-summary mismatch and allocation regret. The theorem term $\mathbb{E}_Z \text{KL}(\pi_Z \| q_Z)$ is defined only after the evaluator predeclares a reproducible construction of the evaluator-side mode summary q_Z . Operationally, the router returns an action, not a mode label or distribution. When no such q_Z construction is frozen, the experiment reports allocation regret as a separate action-frontier diagnostic. After deployment runs, the evaluator estimates the mode-conditional frontier

$$a^*(s) \in \arg \min_{a \in \mathcal{A}} \{\log \kappa(a, s) - \log p(a, s)\},$$

617 or the statistically indistinguishable set of frontier actions. For a selected action $a = \rho_j(x)$ and mode
 618 $s = \sigma(x)$, the corresponding allocation regret diagnostic is

$$r_j(x) = [\log \kappa(a, s) - \log p(a, s)] - \min_{a' \in \mathcal{A}} [\log \kappa(a', s) - \log p(a', s)].$$

619 This regret is computed retrospectively on a diagnostic split and validated on holdout instances; it
 620 is not an oracle input to the router. High information with high allocation regret means that the
 621 progressive signal exposed useful mode structure, but the induced action did not exploit the empirical
 622 cost–success frontier.

623 After these diagnostics are estimated, the experiment still reports the decision-facing first-passage
 624 deployment loss in (5) as final validation. The full-horizon audit loss in (6) is reported separately for
 625 persistence and final-quality checks. These losses are not replacements for the theorem; they test
 626 whether designs favored or explained by the diagnostics improve the practical deployment objective.

627 C Operational decomposition under progressive signals

628 The main four-term identity compares a deployed design to a prior-matched baseline. For the V3
 629 router experiment, it is often useful to compare two progressive signal conditions for the same router.
 630 Assume a finite task mode $S \in \mathcal{S}$ and, for each signal condition j , the conditional mode distribution
 631 induced by the realized signal record

$$\pi_j(s \mid Z_j) = \mathbb{P}[S = s \mid Z_j].$$

632 The fixed orchestrator R induces the routing policy $a_R^j(Z_j) = R(Z_j)$. If an evaluator-side mode
 633 summary $q_R^j(\cdot \mid Z_j) \in \Delta(\mathcal{S})$ has been predeclared, define

$$\mathcal{E}_R^j(s, Z_j) = \frac{\kappa(a_R^j(Z_j), s)}{p(a_R^j(Z_j), s) q_R^j(s \mid Z_j)}.$$

634 The log-effort gain of signal condition Z_j over Z_0 is

$$\Delta_R^{\log}(j) := \mathbb{E}[\log \mathcal{E}_R^0(S, Z_0) - \log \mathcal{E}_R^j(S, Z_j)].$$

635 Expanding the cross-entropy terms gives

$$\Delta_R^{\log}(j) = C_R(j) + \Phi_R(j) + G(j) + M_R(j),$$

636 where

$$\begin{aligned} C_R(j) &= \mathbb{E} \left[\log \frac{\kappa(a_R^0(Z_0), S)}{\kappa(a_R^j(Z_j), S)} \right], \\ \Phi_R(j) &= \mathbb{E} \left[\log \frac{p(a_R^j(Z_j), S)}{p(a_R^0(Z_0), S)} \right], \\ G(j) &= H(S \mid Z_0) - H(S \mid Z_j), \\ M_R(j) &= \mathbb{E} \left[\text{KL}(\pi_0(\cdot \mid Z_0) \parallel q_R^0(\cdot \mid Z_0)) - \text{KL}(\pi_j(\cdot \mid Z_j) \parallel q_R^j(\cdot \mid Z_j)) \right]. \end{aligned}$$

637 When Z_j refines Z_0 , $G(j) = I(S; Z_j \mid Z_0)$. This signal-mode decomposition is theory-facing and
 638 excludes measurement cost unless such a term is added explicitly. It is a diagnostic companion to the
 639 deployment-gain estimand (8), not a replacement for it.

640 D Operational estimator details

641 **Trajectory record.** Each run stores the worker alias, resolved model identifier, API date, split,
 642 task family, task mode, seed, maximum horizon H , signal condition, prompt hash, router model
 643 identifier, router temperature/top-p, router seed policy, parser version, router output, selected worker,
 644 and cost convention. Each step stores the prompt-visible state, proposed `train.py` artifact, parse
 645 status, checker result, validation loss, validation accuracy, relative improvement, incremental worker
 646 wall time, token usage, and checker runtime for audit. Parser failures are logged as unsuccessful
 647 routing or deployment records under the predeclared parser-failure policy.

Table 5: Frozen operational configuration for the main protocol. All values are fixed within a task mode; repeated runs vary only the agentic trajectory.

Component	Fixed value
Dataset and split	CIFAR-10, 50,000 training examples and 10,000 validation examples; balanced subset, label noise 0.0, imbalance ratio 1.0; the train/validation split is fixed in the main protocol
Data seed	Common fixed data/subset seed; alternative split seeds are reported only as robustness variants
<code>train.py</code> initialization seed	SEED=42; used for Python, NumPy, PyTorch, CUDA seeds, deterministic PyTorch mode, and the training-loader generator; alternative initialization seeds are robustness variants
Training transforms	random crop with padding 4, random horizontal flip, CIFAR-10 normalization; validation uses normalization only
Training budget per checker call	fixed-step mode with <code>A_MAX_STEPS</code> =256; time budget 120s is a fallback, not the primary scoring convention
Common hyperparameters	<code>BATCH_SIZE</code> =64, <code>NUM_WORKERS</code> =0, Adam, learning rate $5 \cdot 10^{-4}$, weight decay 10^{-4} , betas (0.9, 0.999), no LR schedule
Task mode	<code>mlp_flat</code> , <code>cnn_compact</code> , or <code>resnet_micro</code>
Shared template parameters	depth 2, base channels 12, channel multiplier 2, batch norm enabled, dropout 0.0, and hidden width 48
Agent horizon	$H = 20$ edit-verify steps; trajectories continue to H after first success to measure persistence and final quality
Progress normalization	$\varepsilon_L = 10^{-8}$ in the denominator of G_t
Primary loss weights	$\alpha_{\text{fail}} = 1$; normalized worker wall-clock is the primary scalar C_δ ; token cost is a sensitivity analysis
Audit loss weights	$\alpha_{\text{occ}} = \alpha_{\text{qual}} = 1$, $\psi(u) = 1 - u$
Measurement cost	pre-deployment probes are charged through ℓ_{meas} using the same wall-clock normalizer as C_δ ; context-token cost is reported as sensitivity
Router configuration	deterministic prompt-only router with fixed model identifier, prompt hash, parser, temperature, top-p, and seed policy recorded before evaluation

648 **Success and time-to-verification.** For trajectory i , τ_i is the first step at which the relative improve-
649 ment threshold is crossed; censored failures have $\tau_i = \infty$. For each action-mode cell,

$$\hat{F}_a(s, h) = \frac{1}{n_{a,s}} \sum_i \mathbf{1}\{\tau_i \leq h\}, \quad \hat{S}_a(s, h) = 1 - \hat{F}_a(s, h).$$

650 The empirical hazard is

$$\hat{\lambda}_a(s, h) = \frac{\sum_i \mathbf{1}\{\tau_i = h\}}{\sum_i \mathbf{1}\{\tau_i \geq h\}}.$$

651 The raw first-success estimate is $\hat{p}(a, s) = \hat{F}_a(s, H)$. For log-effort scores, the protocol uses the
652 predeclared smoothed estimate

$$\tilde{p}(a, s) = \frac{k_{a,s} + 1/2}{n_{a,s} + 1},$$

653 where $k_{a,s}$ is the number of first successes in the cell. This prevents zero-success pilot cells from
654 making log scores infinite; raw success counts are still reported.

655 **Cost.** The primary cost is cumulative worker wall-clock time to $\tau \wedge H$. Secondary worker-side
656 cost is cumulative token count. Checker runtime is logged as a protocol audit field and excluded from
657 action-cost comparisons unless the checker protocol differs across actions. All plots that support
658 model recommendations must be reproducible under at least worker wall-clock and token-count
659 conventions. When a worker fails before producing a parseable artifact, its cost is still counted and
660 the step is marked unsuccessful. The scalar estimator matched to Assumption 1 is the cell mean

$$\hat{\kappa}(a, s) = \frac{1}{n_{a,s}} \sum_{i: S_i = s} C_{\delta,i}(a),$$

661 or its fixed-prototype analogue $\hat{\kappa}_{\chi_s}(a)$ in Section 5.




Figure 8: First-passage diagnostics. Left: first-hit CDFs by mode. Right: cost-to-threshold by mode and worker. These are the direct empirical objects behind the retry and time-to-certified-solution interpretation.




Figure 9: Worker-side cost and accounting diagnostics. These plots check whether conclusions are driven by wall-clock cost, token composition, or quality differences.

662 **Information, allocation summaries, and regret.** For each progressive signal Z_j , a Shannon
663 information estimate is reported only if the protocol predeclares a finite signal representation or
664 calibrated predictive model for estimating $\pi_j(\cdot | Z_j)$ on held-out instances. Otherwise the report
665 uses qualitative signal summaries and router-response diagnostics without calling them estimates
666 of $I(S; Z_j)$. The KL mode-summary divergence term is reported only if the evaluator also freezes
667 a reproducible construction of $\hat{q}_R^j(\cdot | Z_{j,i})$ before seeing evaluation outcomes. In the main fixed-
668 prototype protocol, the primary operational diagnostic is allocation regret,

$$\hat{r}_j(x_i) = [\log \hat{\kappa}(a_{j,i}, s_i) - \log \tilde{p}(a_{j,i}, s_i)] - \min_{a \in \mathcal{A}} [\log \hat{\kappa}(a, s_i) - \log \tilde{p}(a, s_i)],$$

669 where $a_{j,i} = R(Z_j(x_i))$ and $s_i = \sigma(x_i)$.

670 **Uncertainty.** Success probabilities use Wilson intervals. Aggregated objectives, paired gains, and
671 calibration diagnostics use bootstrap or randomization tests at the sampling unit declared for the study.
672 If the main protocol remains a fixed-prototype repeated-run panel, paired randomization tests over
673 repeated trajectories are descriptive for that panel and should not be interpreted as population-level
674 inference. If independent task instances are added, uncertainty is reported at the task-instance level.

675 E Additional diagnostic plots

676 This appendix collects auxiliary diagnostics generated from the experimental logs. Some plots include
677 exploratory cells outside the promoted two-worker comparison; the main text uses the Spark-free
678 mature-worker figures in Figures 1 and 2 for the primary comparison.




Figure 10: Router-facing diagnostics. The signal ablation tests whether progressively richer records expose task-mode structure; the router-selection snapshot checks whether the prompt-only router changes action distribution under those records.




Figure 11: Run-coverage diagnostics for router records and holdout confirmation runs. These plots document which analyses are part of the promoted two-worker comparison and which are exploratory.

NeurIPS Paper Checklist

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [Yes]

Justification: The abstract and introduction state the paper’s scope: verifiable agentic systems, a packed latent-mode assumption, decomposition-guided selection, and a controlled fixed-prototype experimental protocol. The discussion section separates theory-facing and deployment-facing claims.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [Yes]

Justification: Section 7 discusses the packed-mode assumption, small live-agent sample sizes, finite design menus, oracle mode labels, and current cost-measurement limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [Yes]

Justification: Assumption 1 states the assumptions for the main theoretical result, Theorem 1 includes a proof sketch, and Appendix A gives the algebraic proof and pilot concentration details.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

704 Answer: [Yes]
705 Justification: Section 5 describes the AutoResearch task instantiation, task modes, model
706 menu, router-visible signals, horizons, metrics, and evaluation protocol; Appendix D speci-
707 fies the required logs, estimators, and run-metadata fields.

708 **5. Open access to data and code**
709 Question: Does the paper provide open access to the data and code, with sufficient instruc-
710 tions to faithfully reproduce the main experimental results, as described in supplemental
711 material?
712 Answer: [No]
713 Justification: The paper is written against an existing repository and includes reproducibility
714 commands, but an anonymized public code/data release is not included in the current
715 submission package.

716 **6. Experimental setting/details**
717 Question: Does the paper specify all the training and test details (e.g., data splits, hyperpa-
718 rameters, how they were chosen, type of optimizer) necessary to understand the results?
719 Answer: [Yes]
720 Justification: Section 5 describes the AutoResearch task instantiation, task modes, model
721 menu, router-visible signals, primary estimators, horizons, and evaluation design. Additional
722 implementation details are listed in Appendix D.

723 **7. Experiment statistical significance**
724 Question: Does the paper report error bars suitably and correctly defined or other appropriate
725 information about the statistical significance of the experiments?
726 Answer: [Yes]
727 Justification: Section 5 and Appendix D specify Wilson intervals, bootstrap intervals,
728 survival curves, calibration diagnostics, and held-out evaluation for the reported protocol.

729 **8. Experiments compute resources**
730 Question: For each experiment, does the paper provide sufficient information on the com-
731 puter resources (type of compute workers, memory, time of execution) needed to reproduce
732 the experiments?
733 Answer: [No]
734 Justification: The paper reports wall-clock costs and run counts but does not provide a
735 complete compute-resource table for every experiment.

736 **9. Code of ethics**
737 Question: Does the research conducted in the paper conform, in every respect, with the
738 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines>?
739 Answer: [Yes]
740 Justification: The work studies verifiable benchmark tasks and does not involve human
741 subjects, sensitive personal data, or released high-risk models. The submission preserves
742 anonymity and follows NeurIPS code and data policies.

743 **10. Broader impacts**
744 Question: Does the paper discuss both potential positive societal impacts and negative
745 societal impacts of the work performed?
746 Answer: [Yes]
747 Justification: Section 7 notes positive uses in cost-aware and transparent agent deployment
748 as well as limitations and failure modes. The work is methodological and does not introduce
749 a new generative model or dataset of people.

750 **11. Safeguards**
751 Question: Does the paper describe safeguards that have been put in place for responsible
752 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
753 image generators, or scraped datasets)?

754 Answer: [N/A]
 755 Justification: The paper does not release a pre-trained model, image generator, scraped
 756 dataset, or other asset with high misuse risk. It evaluates agent harnesses on synthetic or
 757 generated benchmark tasks.

758 **12. Licenses for existing assets**

759 Question: Are the creators or original owners of assets (e.g., code, data, models), used in
 760 the paper, properly credited and are the license and terms of use explicitly mentioned and
 761 properly respected?

762 Answer: [Yes]

763 Justification: The related-work section cites existing papers and systems used for position-
 764 ing, and the experimental section describes the benchmark surfaces. Released code and
 765 benchmark assets include explicit license details.

766 **13. New assets**

767 Question: Are new assets introduced in the paper well documented and is the documentation
 768 provided alongside the assets?

769 Answer: [Yes]

770 Justification: The work introduces a code harness and analysis artifacts as research assets,
 771 and Appendix D documents the required logging and configuration fields.

772 **14. Crowdsourcing and research with human subjects**

773 Question: For crowdsourcing experiments and research with human subjects, does the paper
 774 include the full text of instructions given to participants and screenshots, if applicable, as
 775 well as details about compensation (if any)?

776 Answer: [N/A]

777 Justification: The paper does not involve crowdsourcing or research with human subjects.

778 **15. Institutional review board (IRB) approvals or equivalent for research with human**
 779 **subjects**

780 Question: Does the paper describe potential risks incurred by study participants, whether
 781 such risks were disclosed to the subjects, and whether Institutional Review Board (IRB)
 782 approvals (or an equivalent approval/review based on the requirements of your country or
 783 institution) were obtained?

784 Answer: [N/A]

785 Justification: The paper does not involve crowdsourcing or research with human subjects, so
 786 IRB or equivalent approval is not applicable.

787 **16. Declaration of LLM usage**

788 Question: Does the paper describe the usage of LLMs if it is an important, original, or
 789 non-standard component of the core methods in this research? Note that if the LLM is used
 790 only for writing, editing, or formatting purposes and does *not* impact the core methodology,
 791 scientific rigor, or originality of the research, declaration is not required.

792 Answer: [Yes]

793 Justification: LLMs are central to the evaluated agent designs. Section 5 describes the
 794 three-worker menu, fixed prompt-only router, sequential signal protocol, model-identifier
 795 logging, and routing-output requirements.